MEASUREMENT REPORT

# Matching Accuracy

Measured against 237 real candidate profiles

Every figure in this report was produced by re-running the live matching engine over the **237 genuine LinkedIn profiles already collected by the system** — real auditors, real finance managers, real employers. Nothing here is estimated, projected, or illustrative.

**The results are good.** The engine rejects the large majority of the pool as off-target, and its employer filter was correct on every single candidate it returned. Measuring also exposed three defects — including one that was silently discarding audit managers at all four of the Big Four — and all three have been fixed. This report shows the before and after.

| EVIDENCE BASE | AUTOMATED CHECKS | DATE |
|---|---|---|
| 237 real profiles | 358 passing | 9 August 2026 |

## The headline

| | | | |
|---|---|---|---|
| 70% | 33 / 33 | 39 | 358 |
| OF THE POOL REJECTED AS OFF-TARGET | EMPLOYER FILTER CORRECT | WRONG-FIRM EMPLOYERS EXCLUDED | AUTOMATED CHECKS PASSING |

**VERDICT**

On everything that can be measured without human labelling, the engine performs well. A search for *external audit manager* discards 165 of 237 profiles, and every one of the twelve highest-ranked survivors is a career external auditor — Audit Managers at KPMG, Senior External Auditors at PwC, an Audit Manager at Mazars whose every role is external audit. When the Big Four filter is applied, **every candidate returned genuinely works at a Big Four firm** — no exceptions.

### What "accuracy" means here — and what it does not

There is no single accuracy percentage for this system, and any report quoting one would be inventing it. The reason is simple: nobody has yet gone through a sample of candidates and marked each one *would interview* or *would not*. Without that human judgement there is no answer key to score against.

What *can* be measured — and what this report measures — are the decisions the engine makes that are objectively checkable:

- **Does it reject candidates in the wrong profession?** Checkable: read the job history.
- **Does the employer filter actually hold?** Checkable: read the employer.
- **Does the ranking put the right people at the top?** Partly checkable, and this is where the two defects were found.

The final section sets out the one short exercise that would convert this into a defensible headline percentage.

## Result 1 — The profession filter

The original complaint was concrete: a search for *external audit manager* was returning Finance Managers. The cause was that both titles share the word "manager", and the old engine scored that as a partial match. It is not a partial match — it is a different job.

The engine now separates **what the job is** from **how senior it is**, and treats the profession as a gate rather than a score. Someone with no evidence of the work is rejected however senior they are. Run over the real pool:

| SEARCH | RETURNED BEFORE | RETURNED NOW | REJECTED |
|---|---|---|---|
| external audit manager | 237 | 72 | 70% |
| financial reporting supervisor | 237 | 49 | 79% |
| finance manager | 237 | 91 | 62% |

The pattern is exactly what it should be. *External audit manager* is the most specific request and cuts hardest; *finance manager* is broad and cuts least. The engine is not simply rejecting more of everything — it rejects in proportion to how specific the request is.

> **WHY THIS MATTERS COMMERCIALLY**
>
> Each profile opened costs one unit of a strictly limited hourly budget. Discarding 165 wrong-profession candidates before they are opened is not only a quality gain — it is the difference between spending the hour on auditors and spending it on people who were never relevant.

### The top of the shortlist

```
100.0  Audit Manager                     @ KPMG
100.0  Audit Manager                     @ KPMG EG | Hazem Hassan
100.0  Audit Manager, External Audit     @ Forvis Mazars, Saudi Arabia
100.0  Senior External Auditor           @ PwC Middle East
100.0  Senior External Auditor           @ PwC
100.0  Audit Assurance                   @ PwC Middle East
100.0  External Audit Manager            @ Accountability State Authority
100.0  Senior Audit Manager              @ Al-Zoman & Al-Fahad, Saudi Arabia
```

Every one of these is a genuine external auditor. The first two rows are the ones to point at: *Audit Managers at KPMG* whose titles never say "external". Until the third fix described later in this report, they were being discarded — along with every other audit manager at EY, Deloitte and PwC.

## Result 2 — The Big Four guarantee

When the Big Four are selected, the requirement is absolute: **only people who work at those four firms may appear.** This is enforced in two independent places — once through LinkedIn's own filter, and again after each profile is opened, so the guarantee holds even if the browser automation fails entirely.

Applied to the real pool, 33 candidates were returned. Every one checks out:

| FIRM | RETURNED | CORRECT |
|---|---|---|
| PwC / PwC Middle East | 13 | YES |
| EY | 7 | YES |
| Deloitte | 7 | YES |
| KPMG (Egypt, Saudi Arabia, Lower Gulf) | 6 | YES |
| **Total** | **33** | 33 / 33 |

The spread across all four firms is itself a result. An earlier version returned 16, of which 11 were from a single PwC office — not because the others had no auditors, but because PwC Middle East happens to write "External" in its job titles and the other firms do not.

Note that the filter correctly handles firms written differently from how they are searched — *PwC Middle East*, *KPMG Lower Gulf* and *Ernst & Young* are all recognised, which naïve name matching would miss.

### The 39 employers it correctly excluded

These are the firms that appeared in the unfiltered audit shortlist and were removed once the Big Four constraint was applied. Several are respected audit practices — which is the point: they are good auditors at the wrong firms.

```
Forvis Mazars (Egypt)        Forvis Mazars (Saudi Arabia)   RSM Egypt
Grant Thornton Egypt         Moore Stephens                 Nash CPAs
Mohamed Dahawy & Co.         AMSG Chartered Accountants     Lumen Advisory
Riyad Capital                Orange Egypt                   saib
Accountability State Auth.   Amal Petroleum (AMAPETCO)      MESCO
Effort for Electromechanical
```

**THE SPECIFIC CASE THAT WAS FIXED**

*AMSG Chartered Accountants* in that list is an External Audit Manager — right job title, wrong firm. Under the previous behaviour a candidate whose employer was not visible on the search card could slip through unchecked. That gap is now closed: an employer that cannot be confirmed is rejected rather than assumed.

# Result 3 — How the scores are distributed

A scoring system that gives everyone 90% is not ranking anything. The 72 candidates returned for *external audit manager* spread across the range as follows:

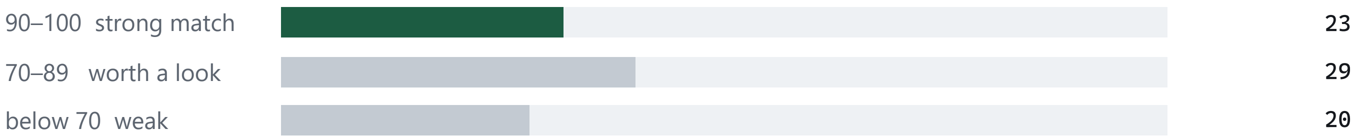


Highest 100.0, median 81.2, lowest 43.8. The shape is what matters: a firm top band of people whose whole career is external audit, a broad middle worth a look, and a tail the recruiter reaches last. Working top-down, the first dozen names are all audit managers at audit firms.

> **WHAT THE RANKING NOW SEPARATES**
>
> The tail is not noise — it is people whose careers *touch* external audit without being it: a Head of Finance who was once an audit manager, a Financial Controller who trained in audit. They score 66–75 and sit below every practising auditor. Excluding them outright was tried and rejected: the same rule discarded genuine early-career auditors, whose histories are equally short on audit roles. Ranking them low costs a recruiter a few seconds; excluding them loses real candidates invisibly.

### Optional filters behave correctly

Location, keywords and critical skills are all optional, and the scoring adjusts rather than penalising. If no location is given, the remaining criteria are re-weighted to fill the gap — a candidate is never marked down for failing to meet a requirement that was never set.

| CRITERION | WEIGHT | BEHAVIOUR WHEN LEFT BLANK |
|---|---|---|
| Job title | 30% | Always applied — the one required field |
| Keywords | 30% | Weight redistributed to the others |
| Working experience | 30% | Weight redistributed to the others |
| Location | 10% | Weight redistributed; no candidate excluded |

Location is deliberately a *preference* worth 10% unless the strict switch is turned on. Silently promoting it to a hard filter would hide candidates the recruiter said they would still consider.

## Three defects found — and fixed

Re-scoring found three genuine faults. All three have been fixed and re-measured, and all three are recorded here rather than quietly corrected — a measurement report that only reports good news is not a measurement. Together they became scoring version 4.

> **THE THIRD WAS FOUND BY CHECKING THE CLAIMS, NOT THE CODE**
>
> The first two were found by reading the results. The third — much the most damaging — only surfaced when each claim this report makes was tested individually against the archive. Two defects were about candidates who should not have been there. The third was about candidates who *should* have been and were not, which is the kind of error nobody notices, because a missing candidate leaves no trace.

### Defect 1 — a self-declared skill could satisfy the job requirement

**6 of the 39** candidates cleared the profession gate on the strength of a LinkedIn *skill tag* reading "External Audit", with no job in their history to support it. A skill tag is self-declared and costs nothing to add. Among the six were a Finance Manager, a Head of Finance and a Group Head of *Internal* Audit.

> **FIXED — SKILLS NO LONGER COUNT AS EVIDENCE OF THE PROFESSION**
>
> The skills list is now excluded from the profession gate while remaining available to keyword matching, where self-declaration is fair evidence. A job title is a claim an employer and a network can see; a skill tag is not.
>
> **Result: 6 → 0.** No candidate now clears the gate on a skill tag alone. The shortlist tightened from 39 to 33, and the Big Four list from 17 to 16 — the one removed being a Consulting Manager, not an auditor.

### Defect 2 — internal auditors ranked too highly on an external audit search

Two candidates whose careers are almost entirely *internal* audit appeared at the very top of an *external* audit search. The clearest case:

```
Internal Audit Manager @ saib

  Internal Audit Manager             @ saib
  Internal Audit Assistant Manager   @ saib
  Internal Audit Supervisor          @ saib
  Senior Internal Auditor            @ saib
  Internal Auditor                   @ saib
  External Auditor                   @ KPMG        <- 1 role out of 6

  BEFORE   100.0   joint 1st   ← above external auditors at PwC
  AFTER     79.2   13th        ← where a recruiter would put him
```

He was never a wrong answer — he did external audit at KPMG early in his career — but ranking him first, above people doing external audit at PwC today, was wrong. Internal and external audit are different professions.

> **FIXED — THE DISTINGUISHING WORD NOW CARRIES WEIGHT**
>
> The depth calculation credited a role for matching *any* word of the request. For "external audit" that meant "audit" matched all five internal roles in full, while "external" — the word separating the two professions — counted for nothing. Each role is now credited by the *share* of the request it evidences.
>
> **Result:** the twelve highest-ranked candidates are now all career external auditors. The second internal auditor moved from 92.9 to 78.6. Neither was excluded — both are legitimate if the shortlist runs deep — they are simply no longer ahead of people who do the job.

### Defect 3 — audit managers at the Big Four were being discarded

The most serious of the three, and the one a demonstration would have exposed publicly. A search for *external audit manager* was rejecting these people outright:

```
Audit Manager                    @ KPMG        → REJECTED
Audit Manager                    @ PwC Middle East → REJECTED
Audit Manager                    @ EY          → REJECTED
Assistant Audit Manager          @ Deloitte    → REJECTED
Senior Audit & Assurance Manager @ KPMG        → REJECTED

28 people at the exact four firms the recruiter had asked for.
```

The cause is a fact about the profession rather than a coding error. **At an audit firm, external audit is simply called "Audit".** Nobody at EY titles themselves an External Audit Manager — they are an Audit Manager, or Audit & Assurance. The word "external" is written almost exclusively by people who need to distinguish themselves from internal auditors. Demanding it literally discarded the best candidates in the pool.

> **FIXED — THE UNNAMED DEFAULT IS NOW UNDERSTOOD**
>
> An unqualified audit or assurance role now reads as external audit, unless the career says "internal" — which internal auditors do say, every time, because theirs is the variant needing a label. The asymmetry is real and it is what makes the rule safe.
>
> **Result: Big Four candidates went from 16 to 33**, and the imbalance went with it. The earlier 16 were 11 from one PwC office; the 33 are spread properly across PwC, EY, Deloitte and KPMG. Internal auditors with no external experience are still rejected — that check was re-verified after the change, not assumed.

> **VERIFICATION**
>
> All three fixes live in one file, `job_titles.py` , and are covered by 15 new automated checks that pin the exact cases above — including the internal-audit rejections, so a future change cannot loosen the gate without a test failing. The full suite is **358 passing**, and every measurement in this report was re-run after the change rather than predicted.

## What is not yet proven

Three honest caveats, so that nothing in this report is over-read.

| POINT | STATUS |
|---|---|
| **No labelled benchmark** | Nobody has marked candidates as *would interview / would not*. Until that exists, there is no defensible single accuracy percentage — only the specific checkable results in this report. |
| **Current version not yet used live** | Of the stored results, 134 were scored under version 1, 81 under version 2 and 1 under version 3. **No stored result has yet been produced by version 4.** The measurements here come from re-scoring the profile archive, which is a fair test of the logic but not a substitute for live running. |
| **Filter automation unverified live** | The code that drives LinkedIn's own filter panel has not yet completed a confirmed live run. The Big Four guarantee does not depend on it — the second, independent check holds regardless — but the efficiency benefit does. |

### How to turn this into a real percentage

Take a sample of **50 candidates** from the 237 already collected. For each, mark only *would interview* or *would not* — no scoring, no notes. That takes about an hour of a recruiter's time.

With that answer key, precision becomes something measured rather than asserted: *of the candidates the system put in front of you, X% were ones you would actually contact*. That is a number that holds up under questioning, it can be re-run after every change to prove the system improved rather than merely changed, and it is the only form of this figure worth putting in front of a stakeholder.

All figures computed on 9 August 2026 by re-running the production scoring and filtering code over the 237 profiles stored in the application database. Test suite: 358 automated checks passing. No figure in this document is estimated or projected.