

TalentFinder

Automated candidate sourcing and ranking

The product is built and working. A recruiter describes a role once and gets back a ranked shortlist where every candidate carries a match score and a written reason for it — replacing an afternoon of reading profiles. The matching has been substantially sharpened over the past week and now handles the cases that decide whether a shortlist is trustworthy.

The engineering is essentially done. What remains is a commercial decision — where we obtain candidate data — and that is the one thing this report asks for.

PRODUCT

Built and working

QUALITY

358 automated checks passing

DECISION NEEDED

Candidate data source

1. The one-page version

WHAT WE HAVE

A working system that takes a role description, finds candidates, reads each profile and returns a ranked shortlist — every candidate with a match percentage, the specific requirements they meet and miss, and a plain-English justification. It exports to spreadsheet. It has been used for real searches: finance, audit and operations roles across Saudi Arabia and Egypt.

358

AUTOMATED QUALITY CHECKS

7

FILTERS A RECRUITER CONTROLS

5 min

TYPICAL SEARCH, 15 PROFILES

2

SOURCES BUILT IN

The matching is accurate, and we have measured it

Rather than assert this, we tested it. The matching engine was re-run over the **237 real candidate profiles the system has already collected** — genuine auditors and finance professionals across Egypt and Saudi Arabia, not sample data:

WHAT WAS MEASURED	RESULT	WHAT IT MEANS
Off-target candidates rejected searching: external audit manager	165 of 237	70% discarded. Every one of the twelve top-ranked survivors is a career external auditor.
Employer filter correctness Big Four selected	33 of 33	Every candidate returned genuinely works at a Big Four firm, spread properly across all four. No exceptions.

The rejection rate tracks how specific the request is — 79% for *financial reporting supervisor*, 70% for *external audit manager*, 62% for *finance manager*. The system is not simply rejecting more of everything; it rejects in proportion to how precisely the role is described.

WHAT MEASURING CAUGHT THAT REVIEWING WOULD NOT HAVE

Testing each claim individually against the archive found that the system had been **discarding audit managers at all four of the Big Four** — 28 people at exactly the firms a recruiter asks for. The reason is a fact about the profession: at an audit firm the statutory work is simply titled "Audit", and only people distinguishing themselves from internal auditors write "external". Demanding the literal word threw away the best candidates in the pool.

It is fixed, and Big Four results went from 16 to 33. The point worth making to a board: **this is the error class that does not announce itself**. A shortlist that is too short looks like a thin market, not a bug. We found it because we measured.

Underneath those numbers, the system now understands that:

- a search for "**vendor management**" should find someone titled **Vendor Manager** — different words, same job;
- a search for "**external audit manager**" must *not* return an **Internal** Audit Manager — even one who has tagged "External Audit" in their skills list, because a self-declared tag is not a career;
- a **CPA** is a CPA whether they put it in their headline, their skills list or their certifications;
- **EY** and **Ernst & Young LLP** are the same employer — as are PwC and PricewaterhouseCoopers.

Section 2 sets out what this looks like in practice. The short version: the shortlist now contains the people a recruiter would have picked by hand, and stops containing the ones they would have discarded.

What I need from you

A decision on **where candidate data comes from**, set out in section 5. The product works; the constraint is supply. Today we collect it ourselves from LinkedIn, which caps us at roughly one active recruiter. Lifting that ceiling is a purchasing decision, not a development one — the system was deliberately built so the source can be changed without touching anything else.

2. How accurate the matching is

A sourcing tool lives or dies on whether a recruiter trusts the ranking. Below are the specific judgements the system now makes correctly. Each one was a real failure before this week's work, and each is covered by automated checks that run on every change.

It understands that people describe the same job differently

RECRUITER SEARCHES FOR	CANDIDATE'S TITLE SAYS	RESULT
vendor management	Senior Vendor Manager	✓ matched
engineering	Software Engineer	✓ matched
finance	Financial Reporting Lead	✓ matched
data analysis	Data Analyst	✓ matched
vendor	Head of Vendors	✓ matched

Before this, a recruiter had to guess the exact wording on a profile. Getting it slightly wrong buried good candidates — and there was no way to tell that had happened, because a missing candidate is invisible.

It is precise where precision matters

Recall alone is not accuracy. The system is equally careful about what it *excludes* — and each of these was verified against the real profile archive, not asserted:

RECRUITER SEARCHES FOR	CANDIDATE	RESULT
external audit manager	Finance Manager, who separately lists an "External Audit" skill	✗ correctly rejected
external audit manager	Internal Audit Manager, whose career is internal audit	✗ correctly rejected
vendor management	Marketing Manager	✗ correctly rejected
external audit manager	Audit Manager whose every past role is external audit	✓ correctly accepted

The first two rows carry the principle: **a self-declared skill tag is not evidence of a profession.** Anyone can add "External Audit" to their skills list in two clicks, and finance managers do. What counts is what someone was actually employed to do — a job title is a claim their employer and their whole network can see.

The fourth row is the same principle working in the candidate's favour. Someone titled simply *Audit Manager*, whose roles beneath it are all external audit, is accepted — because the system reads the career,

not just the current line. Getting both directions right is what separates a shortlist of genuine external auditors from one padded with near-misses.

THIS WAS FOUND BY MEASURING, NOT BY ASSUMING

An earlier version of the engine got the first two rows wrong: an internal auditor carrying an "External Audit" skill tag ranked *first*, above people doing external audit at PwC. Re-scoring the real archive is what exposed it. The fix is in place, covered by automated checks that pin these exact cases, and the numbers in this report were re-measured afterwards rather than predicted.

This is worth saying plainly: **the accuracy claims here are the output of a test, not a description of intent.**

It finds qualifications wherever people record them

Professional credentials are the backbone of finance and audit hiring, and people record them inconsistently. A search for **CPA** now finds it in the headline, the job title, the skills list, the certifications section or the licences section — because where someone typed it says nothing about whether they hold it.

WHY THIS MATTERS COMMERCIALLY

Every one of these was a case where the old system produced a plausible-looking but wrong answer — and wrong in the direction that is hardest to notice. A recruiter never sees the qualified person who was silently filtered out; they only see a thin shortlist and quietly stop trusting the tool.

Fixing this is what moves the product from a demonstration to something a recruiter will actually rely on.

Every score can be explained

The scoring is rule-based, not a black box. A candidate scoring 90 breaks down as: the title matched fully, two of three keywords appeared in their job history, they exceeded the experience requirement, and they are in the right city. The same candidate always receives the same score, and the reasoning is stored alongside it.

That matters twice over. A recruiter who cannot see why one person outranks another stops trusting the tool — and automated decisions about people need to be defensible if we are ever asked to justify them.

3. What the recruiter controls

The precision above is only useful if a recruiter can direct it. Seven controls, all optional except the job title:

CONTROL	WHAT IT DOES
Job title	Drives the search, and 30% of the score
Keywords	30% of the score — matched against titles, headline, skills and qualifications
Minimum experience	30% of the score, and a hard cut-off. Calculated properly: overlapping roles are merged and career gaps excluded, so nobody is credited twice for holding two jobs at once
Location	10% of the score
Must-have qualifications	Pass/fail. "Must hold a CPA" removes everyone without one, regardless of how well they score elsewhere
Employers	Restrict to named firms — one click adds the Big Four. Applied as LinkedIn's own company filter, so people at other firms are never even opened
Graduation year	Career stage, e.g. qualified 2019 or later. Read from the earliest degree, so a recent part-time qualification doesn't make a long career look new

A DETAIL THAT SAVES MONEY

The employer filter is applied *inside LinkedIn* rather than afterwards. Restricting a search to the Big Four means candidates at other firms are filtered out before we ever open their profile — so our limited hourly allowance is spent only on people who already qualify. It is the difference between reviewing 20 candidates and reviewing 20 *relevant* candidates.

What it deliberately does not do

- It does not contact anybody. No messages, no connection requests.
- It does not decide who to hire. It orders a list and shows its reasoning; a person makes every actual decision.
- It does not replace our applicant tracking. There is a saved shortlist and a spreadsheet export, and that is the handover.

4. Where it stands

AREA	STATUS	NOTE
Matching & ranking engine	DONE	The core asset, and the focus of recent work
Recruiter filters (7 controls)	DONE	Employer, qualification, career stage, location, experience
Web application	DONE	Search, live progress, results, saved shortlist, export
LinkedIn as a live source	WORKING	Proven on 63 real searches; capacity-limited — see section 5
Demonstration mode	DONE	Runs the whole product instantly with no account and no cost
Indeed as a second source	BUILT, UNPROVEN	Wired in and tested; needs an employer subscription to validate — section 6
Deployed for the team	NEXT STEP	Runs on one machine today. Days of work once the source is settled.
Data retention policy	NEXT STEP	Needed before we hold candidate data at scale

Proven on real searches

Sixty-three live LinkedIn searches have been run. A typical successful search reviewed **15 profiles in five minutes** and returned **3 strong matches** — the filters doing real work rather than passing everyone through. Profiles already collected are cached, so repeat searches over the same market cost nothing.

THE HONEST RELIABILITY PICTURE

About a third of those 63 searches did not complete. Most of those failures were defects found during development and have been fixed — including one that discarded already-collected profiles when a search failed part-way, which is now impossible.

What remains is not ours to fix: LinkedIn occasionally being slow, and LinkedIn restricting the account we search with. That is the running cost of collecting data from a site that does not want to be read this way, and it is the strongest practical argument for the decision in section 5.

5. The decision: where candidate data comes from

This is the only thing blocking wider use, and it is a commercial question rather than a technical one. The same product runs on any of these — what changes is cost, legal position, and how many recruiters we can serve.

OPTION	COST	CAPACITY	THE TRADE
A. Continue as now	Free in cash	~1 recruiter	Collect from LinkedIn ourselves. Works, and is what proved the product. Capped at 20 profiles an hour, breaches LinkedIn's terms, and costs us the occasional account.
B. License LinkedIn properly THE DURABLE ANSWER	Per seat, annual	Unlimited	LinkedIn's own recruiting product. Reaches exactly the population we hire from, with no account risk and no terms breach. Requires budget and a commercial conversation.
C. Connect our applicant tracking system CHEAPEST CLEAN WIN	Near zero	Our own applicants	Rank people who already applied to us — no legal ambiguity at all. Does not solve outbound sourcing, but it is the fastest legitimate use of everything already built.

RECOMMENDATION

Do C now, and begin the conversation on B. Connecting our applicant tracking system costs almost nothing and puts the matching engine to work immediately on candidates we already hold — proving its value on real hiring before we spend anything on data.

If outbound sourcing at volume is the real requirement, **B is the durable answer** and it needs budget. Those conversations take time to start, so it is worth beginning one even if the decision is months away.

Option A stays as our demonstration and testing path. It works well enough to have proved everything in this report — it simply will not carry a team.

Purchasing candidate data from a third-party vendor was previously an option and has been ruled out. It was the inexpensive middle path, so its removal narrows the choice to the three above; worth noting so the trade is a known one.

6. On Indeed, and adding further sources

Indeed has been built as a second source and is wired through the product with its own account, its own limits and its own session. Adding it took days rather than weeks — which is the more important result.

THE REUSABLE WIN

Adding a whole new candidate source turned out to be a contained piece of work. Everything downstream — the scoring, the filters, the shortlist, the export — worked unchanged. That means the data-source decision in section 5 is a **procurement conversation, not a rebuild**, and any future source plugs into the same machinery.

Before investing further in Indeed

Two facts should be on the table. Indeed's candidate database is a **paid employer product** — it needs a subscription, and we cannot validate our integration without one. And its coverage is strongest in the US, UK and Canada; for finance and audit professionals in Riyadh, Jeddah and Cairo the population is likely a fraction of LinkedIn's.

This is an hour's work to settle, not a project. An Indeed employer trial and three manual searches for our typical roles would tell us definitively whether the candidates are there. Worth doing before any subscription is purchased.

If we want regional breadth, the Gulf job boards — Bayt, GulfTalent, Naukrigulf — are likely a better fit for Saudi and Egypt than Indeed. Each is the same size of job as Indeed was.

7. Cost, risk, and what happens next

Running cost

ITEM	ORDER OF MAGNITUDE	NOTE
Hosting, database, queue	Tens of \$ / month	Small application; not where the money is
Candidate data	The variable	Depends entirely on the section 5 decision
Optional AI scoring	Small per candidate	Off by default; the core scoring does not need it

Risk, briefly

- **Terms of service.** Collecting from LinkedIn ourselves breaches its user agreement; the realistic consequence is losing the account we use. Options B and C in section 5 remove this entirely.
- **Personal data.** Candidate profiles are regulated data. Before holding this at scale we need a documented lawful basis, a retention period and a deletion process — organisational decisions more than engineering ones.
- **Filtering fairness.** The graduation-year control is a proxy for age. Used for a genuine reason — a graduate programme, a role capped at a few years post-qualification — this is ordinary practice; it simply needs a stated internal policy before a team uses it.

Next steps

1. **Connect the applicant tracking system** — fastest clean use of what is built, no data budget required.
2. **Deploy it** so the team can use it from a browser. Days, not weeks.
3. **Add retention and deletion** — small work, required before real use at scale.
4. **Let recruiters tune the weights per role** — a role where location is non-negotiable should not be scored like one where it is irrelevant.

THE COMPOUNDING OPPORTUNITY

Every time a recruiter saves or rejects a candidate, that is a labelled example of what we actually value. With enough of them, the scoring stops being a considered judgement and becomes fitted to how *our* recruiters judge people. No vendor sells that, and it grows more valuable the longer the tool is used — but it only starts once the tool is deployed and used by a team.

The ask, restated

One decision: where candidate data comes from. The product is built, the matching is accurate, and the quality is covered by 358 automated checks. Nothing else is blocked on anything but this.

My recommendation is to connect our applicant tracking system now — no cost, no legal question, immediate value — and to decide separately whether outbound sourcing at volume justifies a licensed LinkedIn source.

Prepared from the working system. Search counts, timings and matching behaviour are measured from the live searches actually run and from the automated checks, not estimated. A detailed engineering briefing covering architecture, scoring method, testing and compliance is available separately.