

TalentFinder

Automated candidate sourcing and ranking

A recruiter describes a role once and gets back a ranked shortlist, every candidate carrying a match score and a written reason for it — replacing an afternoon of reading profiles.

It now works end to end. A real search returned six external audit managers at EY, KPMG and Deloitte, correctly ranked. The accuracy is measured rather than asserted, against 237 real profiles.

The engineering is done. What remains is a commercial decision — where candidate data comes from — and that is what this report asks for.

PRODUCT

Working, proven on a live search

QUALITY

419 automated checks passing

DECISION NEEDED

Where the data comes from

1. Where it stands

THE PITCH

Recruiters spend most of their sourcing time reading profiles that were never going to fit. TalentFinder automates that first pass. You describe the role — title, qualifications, minimum experience, employers — and the system finds candidate profiles, reads each one, and returns a ranked shortlist where every candidate carries a match percentage, a breakdown showing *why*, and the specific things they have and lack. It exports to spreadsheet.

6 / 6

LIVE SEARCH: CORRECT CANDIDATES

70%

OFF-TARGET PROFILES REJECTED

33 / 33

EMPLOYER FILTER CORRECT

419

AUTOMATED QUALITY CHECKS

The proof

A recruiter asked for external audit managers at the Big Four, holding a professional qualification, with five years' experience. This is what came back:

100	Hisham Sabry, ACCA, ESAA	Audit Manager	EY
100	Marwan Mohamed, CPA	Audit Manager	KPMG
95	Mohamed Fathy Bakr, ACCA, FESAA	Audit Manager	EY
92	Mohamed Elsherif, ACCA, CA, DipIFR	Audit Manager	EY
92	Belal Mansour, ESAA, DipIFR	Assistant Manager	Deloitte
85	Nancy Hamdy, ACCA	Audit Manager	EY

Six people, all at the requested firms, all genuinely auditors, ranked by how well their career fits the role. That is the product doing its job.

What I need from you

A decision on **where candidate data comes from**, set out in section 4. The product works; the constraint is supply. Today we collect it ourselves from LinkedIn, which caps each recruiter at a fixed hourly rate and puts their personal LinkedIn account at risk. Lifting that is a purchasing decision, not a development one — the system was deliberately built so the source can be changed without touching anything else.

2. How good the matching is

Rather than assert this, we tested it: the engine was re-run over the **237 real candidate profiles the system has collected** — genuine auditors and finance professionals across Egypt and Saudi Arabia.

WHAT WAS MEASURED	RESULT	WHAT IT MEANS
Off-target candidates rejected	70%	The twelve highest-ranked survivors are all career external auditors.
Employer filter correctness	33 / 33	Every candidate returned genuinely worked at a Big Four firm, spread across all four.
Live search	6 / 6	Proven outside the archive, on a real recruiter's search.

The rejection rate tracks how specific the request is — 79% for *financial reporting supervisor*, 70% for *external audit manager*, 62% for *finance manager*. The system is not rejecting more of everything; it rejects in proportion to how precisely the role is described.

It understands this market, not a textbook one

Two things had to be learned from the data, and both would have quietly halved the useful output:

WHAT WE FOUND	WHY IT MATTERED
At an audit firm, external audit is just called "Audit"	Nobody at EY titles themselves "External Audit Manager". Requiring the literal word discarded 28 audit managers at exactly the four firms the recruiter had asked for.
CPA is the American qualification	This market runs on ACCA, ESAA and DiplIFR. Five of the six candidates above hold no CPA. Filtering on it alone would have returned one person out of six.

WHAT MEASURING CAUGHT THAT REVIEWING WOULD NOT HAVE

Both of these were found by testing each claim against the real archive, not by reading the code — which looked correct. The point worth making: **this is the error class that does not announce itself**. A shortlist that is too short looks like a thin market, not a bug. We found them because we measured, and the system now says so when it cannot do what it was asked.

3. Built to be trusted

A sourcing tool lives or dies on whether a recruiter believes the shortlist. The single largest piece of work since the last report was making the system incapable of quietly giving a wrong answer.

SITUATION	WHAT THE RECRUITER IS TOLD
Nobody matched the criteria	Which filter rejected people, and how many — <i>"8 work at a different employer, 2 are missing a required qualification"</i> — so they know what to loosen.
A filter could not be applied	Plainly, before the results: <i>"These results are not what you asked for."</i> Plus how to work around it.
The hourly limit was reached	That the run was cut short, that the criteria were not the problem, and when it can resume.
The LinkedIn account was locked	That it stopped, and that everything collected before it was kept.
Their machine died mid-search	That it stopped unexpectedly, and that re-running will not repeat the cost.

Previously all five showed the same sentence: *"No candidates matched. Try lowering the minimum score."* — which blames the recruiter for a technical failure. That is how a tool loses its users, and it is now impossible.

And the obligations that come with candidate data

A candidate can be erased on request — their profile, every search result referencing them, and every recruiter's saved copy — and the deletion is logged. Profiles and references are removed automatically after six months, and the raw collected payload is discarded once it is no longer needed. This was previously a written promise with no implementation behind it.

4. The decision

Everything above is engineering that is finished. What follows is not, and it is the only thing standing between this and company-wide use.

How data is collected today

Each recruiter's own computer drives a real browser against LinkedIn using **their own personal LinkedIn account**, which they sign into by hand. There is a hard cap on how many profiles an hour that account may open — the control that keeps it alive.

CONSEQUENCE	DETAIL
It breaches LinkedIn's terms	Rolling it out company-wide moves that from one person's account to the company, at company direction.
Accounts get restricted	Two have been, at the current rate. A restriction costs that recruiter their own professional network, not just a tool.
Throughput is capped per person	It scales with headcount rather than with hardware. Five recruiters is five times the rate — but also five accounts at risk.
It breaks when LinkedIn changes	Twice already. Each time it is unplanned work on someone else's schedule.

The alternative

Buying candidate data from a licensed provider removes all four at once: no personal accounts at risk, no terms breach, no per-person cap, and nothing that breaks when a website is redesigned.

Nothing else about the product changes. The data source was built as a replaceable component from the start — the matching, ranking, explanation and export all work identically whatever supplies the profiles. That was a deliberate decision made early precisely so this choice could be made late, commercially, without rewriting the system.

What I am asking for: approval to evaluate licensed providers and a budget range, and a decision on whether recruiters' personal LinkedIn accounts may be used in the meantime.

5. What is not finished

Stated plainly, because you should hear it here first.

ITEM	STATUS
Not yet on a server	It runs on a developer machine today. Deployment is documented and needs a small server and about an hour.
One live search, not a week of them	The pipeline is proven. It is not yet proven to keep working across sustained real use.
No single accuracy percentage	Deliberate. That needs a recruiter to mark a sample of candidates <i>would interview / would not</i> — about an hour — and until then any number quoted would be invented.
One filter needs a manual step	Restricting to specific firms currently works by pasting a link from LinkedIn — ten seconds, entirely reliable. Automating that step is unfinished, and the app says so when it cannot do it.

RECOMMENDATION

Deploy it for a small number of recruiters and let it run for two weeks before any wider rollout. The product is good; what it has not had is time. Two weeks of real use will surface the ordinary problems that only appear with real users, at a scale where they cost patience rather than credibility.

Prepared 14 August 2026. Accuracy figures come from re-running the production code over the 237 candidate profiles in the application database, plus one live LinkedIn search. 419 automated checks passing. Legal statements summarise the project's own compliance documentation and are not legal advice.